



# **AMD Instinct MI355X Kimi K3 Can Generate Up to \$32B of Revenue per GigaWatt per Year White Paper**

Executive Summary – AgentX Serving Economics Analysis

Published by SemiAnalysis InferenceX Team – September 2026

# AMD Instinct MI355X Kimi K3 Can Generate Up to \$32B of Revenue per GigaWatt per Year

Executive Summary – AgentX Serving Economics Analysis · Published by SemiAnalysis InferenceX Team · September 2026

## Summary


InferenceX AgentX measures how many agentic-coding tokens a production serving stack delivers at a fixed interactivity target. Priced at Kimi K3's public OpenRouter rates and scaled to one gigawatt of provisioned power, the MI355X vLLM result is **\$32.6 billion of revenue per year** at P90 45 tok/s/user and 60% utilization.

**Owned at hyperscaler volume (\$1.50 per chip-hour), the fleet keeps \$16.5 billion of profit at a 50.7% margin. Rented on a 3-year contract at August 2026 pricing (\$3.00 per chip-hour), it keeps \$10.3 billion at 31.5%.**

**Figure 1: Kimi K3 2.8T Agentic Revenue & Profit Estimates per GigaWatt per Year at P90 45 tok/s/user**

**Cost Tier:** Owning Hyperscaler · **TCO \$/chip/hr:** MI355X \$1.50 · **Utilization:** 60% · **Model License Fee:** 30% · **Updated:** 2026-09-04 · **Source:** SemiAnalysis InferenceX™

**Selling Price per Million Tokens:** Input \$3 · Cached Input \$0.3 · Output \$15 (OpenRouter)




**Owned at hyperscaler volume.** Compute is charged on every provisioned GPU-hour at \$1.50; the model lab takes 30% of revenue as a license fee; the operator keeps \$16.5B, or \$3.95 per chip-hour.

**Figure 2: Kimi K3 2.8T Agentic Revenue & Profit Estimates per GigaWatt per Year at P90 45 tok/s/user**

**Cost Tier:** 3-Year Rental (August 2026 pricing) · **TCO \$/chip/hr:** MI355X \$3.00 · **Utilization:** 60% · **Model License Fee:** 30% · **Updated:** 2026-09-04 · **Source:** SemiAnalysis InferenceX™

**Selling Price per Million Tokens:** Input \$3 · Cached Input \$0.3 · Output \$15 (OpenRouter)



**Rented at \$3.00 per chip-hour.** Revenue and license fee are unchanged; compute expense doubles to \$12.6B and profit falls to \$10.3B, or \$2.45 per chip-hour. The rented fleet breaks even at 33% utilization; the owned fleet at 16.5%.

## Owned vs. rented, per gigawatt-year

| Cost profile                 | Revenue | Compute expense | Model license fee | Profit         | Margin | Profit / chip-hr |
|------------------------------|---------|-----------------|-------------------|----------------|--------|------------------|
| Owned, \$1.50 per chip-hour  | \$32.6B | \$6.3B          | \$9.8B            | <b>\$16.5B</b> | 50.7%  | \$3.95           |
| Rented, \$3.00 per chip-hour | \$32.6B | \$12.6B         | \$9.8B            | <b>\$10.3B</b> | 31.5%  | \$2.45           |

## How the estimate is built

- **6,115 tok/s per GPU** at P90 45 tok/s/user, interpolated on the measured MI355X vLLM frontier (FP4, MTP, TP8).
- **478,469 GPUs per gigawatt** at 2.09 kW all-in per chip; 60% of benchmarked throughput is sold.
- **\$0.59 per million tokens** blended: 99.2% of tokens are input and 93.7% of those hit the prefix cache.
- **Profit** = revenue – compute (every GPU-hour, busy or idle) – 30% model license fee.

**Sources:** [inferencex.semianalysis.com/profit-estimator-per-gigawatt](https://inferencex.semianalysis.com/profit-estimator-per-gigawatt) · [inferencex.semianalysis.com/agentx](https://inferencex.semianalysis.com/agentx) · [openrouter.ai/moonshotai/kimi-k3](https://openrouter.ai/moonshotai/kimi-k3) · [semianalysis.com/ai-cloud-tco-model](https://semianalysis.com/ai-cloud-tco-model). Figures are estimates from benchmark throughput and public list prices, not forecasts of any operator's results.