




AMD Instinct MI355X Kimi K3 Can Generate Up to \$32B of Revenue per GigaWatt per Year

Executive Summary – Agentic Inference Serving Economics Analysis

Published by SemiAnalysis InferenceX Team – September 2026

AMD Instinct MI355X Kimi K3 Can Generate Up to \$32B of Revenue per GigaWatt per Year

Executive Summary – Agentic Inference Serving Economics Analysis · Published by SemiAnalysis InferenceX Team · September 2026

Summary

InferenceX AgentX measures how many agentic-coding tokens a production serving stack delivers at a fixed interactivity target. Priced at Kimi K3's public OpenRouter rates and scaled to one gigawatt of provisioned power, the MI355X vLLM result is **\$32.6 billion of revenue per year** at P90 45 tok/s/user and 60% utilization.

Owned at hyperscaler volume (\$1.50 per chip-hour), the fleet keeps \$16.5 billion of profit at a 50.7% margin. Rented on a 3-year contract at August 2026 pricing (\$3.00 per chip-hour), it keeps \$10.3 billion at 31.5%.





Figure 1



Owned at hyperscaler volume. Compute is charged on every provisioned GPU-hour at \$1.50; the model lab takes 30% of revenue as a license fee; the operator keeps \$16.5B, or \$3.95 per chip-hour.

\$32.6B

Revenue per GW-year

Kimi K3 at OpenRouter list price,
60% utilization

\$10.3B

Profit per GW-year

Rented at \$3.00 per chip-hour ·
\$2.45 per chip-hour

31.5%

Profit margin

After compute expense and 30%
model license fee

Figure 2




Kimi K3 2.8T Agentic Revenue & Profit Estimates per GigaWatt Per Year at P90 45 tok/s/user Interactivity

Cost Tier: Custom \$/GPU/hr (August 2026 3-Year Rental Pricing) Utilization: 60% Model License Fee Assumption: 30% Updated: 2026-09-04 Source: SemiAnalysis InferenceX™

TCO \$/chip/hr: MI355X: 3 Source: SemiAnalysis Market August 2026 Pricing Surveys & AI Cloud TCO Model ↗

Selling Price per Million Tokens: Input \$3 · Cached Input \$0.3 · Output \$15 (OpenRouter)



Rented at \$3.00 per chip-hour. Revenue and license fee are unchanged; compute expense doubles to \$12.6B and profit falls to \$10.3B, or \$2.45 per chip-hour. The rented fleet breaks even at 33% utilization; the owned fleet at 16.5%.

How the estimate is built

- **6,115 tok/s per GPU** at P90 45 tok/s/user, interpolated on the measured MI355X vLLM frontier (FP4, MTP, TP8).
- **478,469 GPUs per gigawatt** at 2.09 kW all-in per chip; 60% of benchmarked throughput is sold.
- **\$0.59 per million tokens** blended: 99.2% of tokens are input and 93.7% of those hit the prefix cache.
- **Profit** = revenue – compute (every GPU-hour, busy or idle) – 30% model license fee.

Sources: inferencex.semianalysis.com/profit-estimator-per-gigawatt · inferencex.semianalysis.com/agentx-openrouter.ai/moonshotai/kimi-k3 · semianalysis.com/ai-cloud-tco-model. Figures are estimates from benchmark throughput and public list prices, not forecasts of any operator's results.